

wine_visualizations

October 23, 2017

1 Plotting with Matplotlib

Use Matplotlib to create bar charts that visualize the conclusions you made with groupby and query.

```
In [1]: # Import necessary packages and load 'winequality_edited.csv'
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
% matplotlib inline
```

```
df = pd.read_csv('winequality_edited.csv')
df.head()
```

```
Out[1]:
```

	fixed_acidity	volatile_acidity	citric_acid	residual_sugar	chlorides	\
0	7.4	0.70	0.00	1.9	0.076	
1	7.8	0.88	0.00	2.6	0.098	
2	7.8	0.76	0.04	2.3	0.092	
3	11.2	0.28	0.56	1.9	0.075	
4	7.4	0.70	0.00	1.9	0.076	

	free_sulfur_dioxide	total_sulfur_dioxide	density	pH	sulphates	\
0	11.0	34.0	0.9978	3.51	0.56	
1	25.0	67.0	0.9968	3.20	0.68	
2	15.0	54.0	0.9970	3.26	0.65	
3	17.0	60.0	0.9980	3.16	0.58	
4	11.0	34.0	0.9978	3.51	0.56	

	alcohol	quality	color	acidity_levels
0	9.4	5	RED	low
1	9.8	5	RED	med-high
2	9.8	5	RED	med-low
3	9.8	6	RED	med-high
4	9.4	5	RED	low


1.0.1 #1: Do wines with higher alcoholic content receive better ratings?

Create a bar chart with one bar for low alcohol and one bar for high alcohol wine samples. This first one is filled out for you.

```
In [2]: # Use query to select each group and get its mean quality
        median = df['alcohol'].median()
        low = df.query('alcohol < {}'.format(median))
        high = df.query('alcohol >= {}'.format(median))

        mean_quality_low = low['quality'].mean()
        mean_quality_high = high['quality'].mean()

In [3]: # Create a bar chart with proper labels
        locations = [1, 2]
        heights = [mean_quality_low, mean_quality_high]
        labels = ['Low', 'High']
        plt.bar(locations, heights, tick_label=labels)
        plt.title('Average Quality Ratings by Alcohol Content')
        plt.xlabel('Alcohol Content')
        plt.ylabel('Average Quality Rating');
```



1.0.2 #2: Do sweeter wines receive higher ratings?

Create a bar chart with one bar for low residual sugar and one bar for high residual sugar wine samples.

```
In [4]: # Use query to select each group and get its mean quality
sugar_median = df['residual_sugar'].median()
sugar_low = df.query('residual_sugar < {}'.format(sugar_median))
sugar_high = df.query('residual_sugar >= {}'.format(sugar_median))


sugar_mean_quality_low = sugar_low['quality'].mean()
sugar_mean_quality_high = sugar_high['quality'].mean()

print('sugar low quality mean: {}'.format(sugar_mean_quality_low))
print('sugar high quality mean: {}'.format(sugar_mean_quality_high))

sugar low quality mean: 5.808800743724822
sugar high quality mean: 5.82782874617737
```

```
In [5]: # Create a bar chart with proper labels
```

```
#locations = [1, 2]
sugar_heights = [sugar_mean_quality_low, sugar_mean_quality_high]
#labels = ['Low', 'High']
plt.bar(locations, sugar_heights, tick_label=labels)
plt.title('Average Quality Ratings by Residual Sugar Content')
plt.xlabel('Residual Sugar Content')
plt.ylabel('Average Quality Rating');
```



1.0.3 #3: What level of acidity receives the highest average rating?

Create a bar chart with a bar for each of the four acidity levels.

```
In [6]: # Use groupby to get the mean quality for each acidity level
```

```
quality_acid_level_means = df.groupby(['acidity_levels'])['quality'].mean()
print(quality_acid_level_means)
```

```
acidity_levels
high          5.783343
low           5.859415
med-high      5.784540
med-low       5.850832
Name: quality, dtype: float64
```

```
In [7]: # Create a bar chart with proper labels
```

```
# stop hardcoding the locations, we can generate them once we know how many means we
# will plot (i.e. the number of acidity levels we grouped by)
x_coords = list(np.arange(1, len(quality_acid_level_means) + 1))

# reindex because the default order wasn't human-readable friendly
quality_acid_level_means = quality_acid_level_means.reindex(['low', 'med-low', 'med-high', 'high'])
#print(quality_acid_level_means)

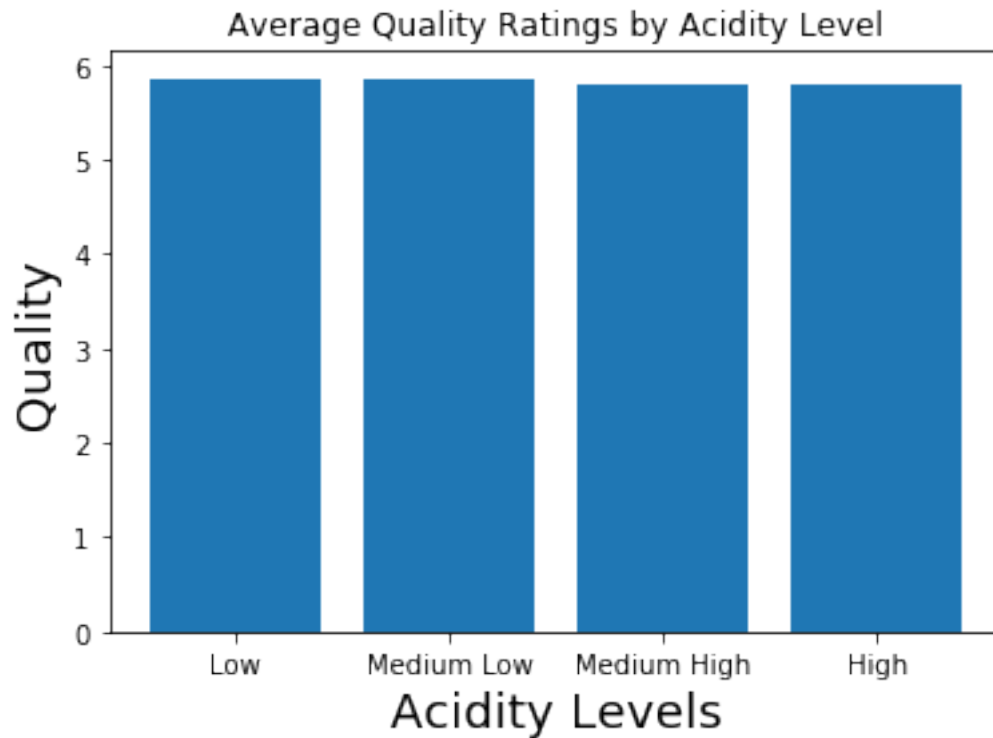
# labels are in same order of reindex
acidity_labels = ['Low', 'Medium Low', 'Medium High', 'High']

plt.bar(x_coords, quality_acid_level_means, tick_label=acidity_labels)

plt.title('Average Quality Ratings by Acidity Level')

plt.xlabel('Acidity Levels', fontsize=18)
plt.ylabel('Quality', fontsize=18)
```

```
Out[7]: Text(0,0.5,'Quality')
```




1.0.4 Bonus: Create a line plot for the data in #3

You can use pyplot's `plot` function for this.

```
In [13]: # this shows the order from the re-index  
quality_acid_level_means
```

```
Out[13]: acidity_levels  
low      5.859415  
med-low  5.850832  
med-high 5.784540  
high     5.783343  
Name: quality, dtype: float64
```

```
In [17]: # need to use xticks to get the labels  
plt.plot(x_coords, quality_acid_level_means, 'bo')  
plt.xticks(x_coords, acidity_labels);
```



Compare this with the bar chart. How might showing this visual instead of the bar chart affect someone's conclusion about this data?