

fix_datatypes_air_pollution

October 26, 2017

1 Fixing air_pollution_score Data Type

- 2008: convert string to float
- 2018: convert int to float

```
In [1]: # load datasets
import pandas as pd
```

```
In [2]: df_08 = pd.read_csv('data_08.csv')
df_08.head(1)
```

```
Out[2]:
```

	model	displ	cyl	trans	drive	fuel	veh_class	\
0	ACURA MDX	3.7	6	Auto-S5	4WD	Gasoline	SUV	

	air_pollution_score	city_mpg	hwy_mpg	cmb_mpg	greenhouse_gas_score	smartway
0		7	15	20	17	4 no

```
In [3]: df_18 = pd.read_csv('data_18.csv')
df_18.head(1)
```

```
Out[3]:
```

	model	displ	cyl	trans	drive	fuel	veh_class	\
0	ACURA RDX	3.5	6	SemiAuto-6	2WD	Gasoline	small SUV	

	air_pollution_score	city_mpg	hwy_mpg	cmb_mpg	greenhouse_gas_score	smartway
0		3	20	28	23	5 No

```
In [4]: df_08['air_pollution_score'].value_counts()
```

```
Out[4]:
```

6	500
7	398
9.5	80
9	7
6/4	1

Name: air_pollution_score, dtype: int64

```
In [6]: df_18['air_pollution_score'].value_counts()
```

```
Out[6]: 3      372
        5      187
        7      139
        1       89
        6        7
        Name: air_pollution_score, dtype: int64
```

```
In [9]: # try using Pandas to_numeric or astype function to convert the
        # 2008 air_pollution_score column to float -- this won't work
        df_08['air_pollution_score'].str.contains('/')
```

```
Out[9]: 0      False
        1      False
        2      False
        3      False
        4      False
        5      False
        6      False
        7      False
        8      False
        9      False
        10     False
        11     False
        12     False
        13     False
        14     False
        15     False
        16     False
        17     False
        18     False
        19     False
        20     False
        21     False
        22     False
        23     False
        24     False
        25     False
        26     False
        27     False
        28     False
        29     False
        ...
        956    False
        957    False
        958    False
        959    False
        960    False
        961    False
```

```

962    False
963    False
964    False
965    False
966    False
967    False
968    False
969    False
970    False
971    False
972    False
973    False
974    False
975    False
976    False
977    False
978    False
979    False
980    False
981    False
982    False
983    False
984    False
985    False
Name: air_pollution_score, Length: 986, dtype: bool

```

2 Figuring out the issue

Looks like this isn't going to be as simple as converting the datatype. According to the error above, the value at row 582 is "6/4" - let's check it out.

```
In [5]: df_08.iloc[582]
```

```

Out[5]: model          MERCEDES-BENZ C300
        displ          3
        cyl            6
        trans          Auto-L7
        drive          2WD
        fuel            ethanol/gas
        veh_class      small car
        air_pollution_score  6/4
        city_mpg        13/18
        hwy_mpg         19/25
        cmb_mpg         15/21
        greenhouse_gas_score  7/6
        smartway        no
Name: 582, dtype: object

```

3 It's not just the air pollution score!

The mpg columns and greenhouse gas scores also seem to have the same problem - maybe that's why these were all saved as strings! According to [this link](#), which I found from the PDF documentation:

"If a vehicle can operate on more than one type of fuel, an estimate is provided for each fuel type."

Ohh.. so all vehicles with more than one fuel type, or hybrids, like the one above (it uses ethanol AND gas) will have a string that holds two values - one for each. This is a little tricky, so I'm going to show you how to do it with the 2008 dataset, and then you'll try it with the 2018 dataset.

```
In [8]: # First, let's get all the hybrids in 2008
```

```
hb_08 = df_08[df_08['fuel'].str.contains('/')]  
hb_08
```

```
Out[8]:
```

	model	displ	cyl	trans	drive	fuel	veh_class	\
582	MERCEDES-BENZ C300	3.0	6	Auto-L7	2WD	ethanol/gas	small car	

	air_pollution_score	city_mpg	hwy_mpg	cmb_mpg	greenhouse_gas_score	smartway
582	6/4	13/18	19/25	15/21		7/6 no

Looks like this dataset only has one! The 2018 has MANY more - but don't worry - the steps I'm taking here will work for that as well!

```
In [10]: # hybrids in 2018
```

```
hb_18 = df_18[df_18['fuel'].str.contains('/')]  
hb_18
```

```
Out[10]:
```

	model	displ	cyl	trans	drive	\
52	BMW 330e	2.0	4	SemiAuto-8	2WD	
78	BMW 530e	2.0	4	SemiAuto-8	2WD	
79	BMW 530e	2.0	4	SemiAuto-8	4WD	
92	BMW 740e	2.0	4	SemiAuto-8	4WD	
189	CHEVROLET Impala	3.6	6	SemiAuto-6	2WD	
195	CHEVROLET Silverado 15	4.3	6	Auto-6	2WD	
196	CHEVROLET Silverado 15	4.3	6	Auto-6	4WD	
197	CHEVROLET Silverado 15	5.3	8	Auto-6	2WD	
212	CHEVROLET Suburban 1500	5.3	8	Auto-6	2WD	
214	CHEVROLET Suburban 1500	5.3	8	Auto-6	4WD	
216	CHEVROLET Tahoe 1500	5.3	8	Auto-6	2WD	
218	CHEVROLET Tahoe 1500	5.3	8	Auto-6	4WD	
225	CHEVROLET Volt	1.5	4	CVT	2WD	
226	CHEVROLET Volt	1.5	4	CVT	2WD	
227	CHRYSLER 300	3.6	6	Auto-8	2WD	
229	CHRYSLER 300	3.6	6	Auto-8	4WD	
244	DODGE Charger	3.6	6	Auto-8	2WD	
246	DODGE Charger	3.6	6	Auto-8	4WD	

300	FORD Fusion Energi Plug-in Hybrid	2.0	4	CVT	2WD
326	GMC Sierra 15	4.3	6	Auto-6	2WD
327	GMC Sierra 15	4.3	6	Auto-6	4WD
328	GMC Sierra 15	5.3	8	Auto-6	2WD
345	GMC Yukon 1500	5.3	8	Auto-6	2WD
347	GMC Yukon 1500	5.3	8	Auto-6	4WD
351	GMC Yukon 1500 XL	5.3	8	Auto-6	2WD
354	GMC Yukon XL 1500	5.3	8	Auto-6	4WD
442	JEEP Cherokee	2.4	4	Auto-9	2WD
444	JEEP Cherokee	2.4	4	Auto-9	4WD
462	KARMA Revero	2.0	4	Auto-1	2WD
571	MERCEDES-BENZ CLA250 4Matic	2.0	4	AutoMan-7	4WD
578	MERCEDES-BENZ GLA250 4Matic	2.0	4	AutoMan-7	4WD
584	MERCEDES-BENZ GLE350 4Matic	3.5	6	Auto-7	4WD
616	MINI Cooper SE Countryman All4	1.5	3	SemiAuto-6	4WD
742	TOYOTA Sequoia FFV	5.7	8	SemiAuto-6	4WD
747	TOYOTA Tundra FFV	5.7	8	SemiAuto-6	4WD
777	VOLVO S90	2.0	4	SemiAuto-8	4WD
789	VOLVO XC 60	2.0	4	SemiAuto-8	4WD
793	VOLVO XC 90	2.0	4	SemiAuto-8	4WD

	fuel	veh_class	air_pollution_score	city_mpg	hwy_mpg	\
52	Gasoline/Electricity	small car	3	28/66	34/78	
78	Gasoline/Electricity	small car	7	27/70	31/75	
79	Gasoline/Electricity	small car	7	27/66	31/68	
92	Gasoline/Electricity	large car	3	25/62	29/68	
189	Ethanol/Gas	large car	5	14/18	20/28	
195	Ethanol/Gas	pickup	5	12/18	16/24	
196	Ethanol/Gas	pickup	5	12/17	15/22	
197	Ethanol/Gas	pickup	3	12/16	17/23	
212	Ethanol/Gas	standard SUV	3	12/16	17/23	
214	Ethanol/Gas	standard SUV	3	11/16	15/22	
216	Ethanol/Gas	standard SUV	3	12/16	17/23	
218	Ethanol/Gas	standard SUV	3	11/16	16/22	
225	Gasoline/Electricity	small car	3	43/113	42/99	
226	Gasoline/Electricity	small car	7	43/113	42/99	
227	Ethanol/Gas	large car	3	14/19	22/30	
229	Ethanol/Gas	large car	3	13/18	20/27	
244	Ethanol/Gas	large car	3	14/19	22/30	
246	Ethanol/Gas	large car	3	13/18	20/27	
300	Gasoline/Electricity	midsize car	7	43/102	41/91	
326	Ethanol/Gas	pickup	5	12/18	16/24	
327	Ethanol/Gas	pickup	5	12/17	15/22	
328	Ethanol/Gas	pickup	3	12/16	17/23	
345	Ethanol/Gas	standard SUV	3	12/16	17/23	
347	Ethanol/Gas	standard SUV	3	11/16	16/22	
351	Ethanol/Gas	standard SUV	3	12/16	17/23	
354	Ethanol/Gas	standard SUV	3	11/16	15/22	

442	Ethanol/Gas	small SUV	3	15/21	22/30
444	Ethanol/Gas	small SUV	3	14/21	21/28
462	Gasoline/Electricity	small car	1	20/59	21/61
571	Ethanol/Gas	small car	5	17/24	24/32
578	Ethanol/Gas	small SUV	5	17/23	23/31
584	Ethanol/Gas	standard SUV	3	13/18	17/22
616	Gasoline/Electricity	midsize car	3	28/63	27/66
742	Ethanol/Gas	standard SUV	5	9/13	13/17
747	Ethanol/Gas	pickup	5	9/13	12/17
777	Gasoline/Electricity	midsize car	7	26/70	33/72
789	Gasoline/Electricity	small SUV	7	26/60	28/58
793	Gasoline/Electricity	standard SUV	7	26/63	30/61

	cmb_mpg	greenhouse_gas_score	smartway
52	30/71	10	Yes
78	29/72	10	Elite
79	28/67	10	Elite
92	27/64	9	Yes
189	16/22	4	No
195	14/20	4	No
196	13/19	3	No
197	14/19	3	No
212	14/19	3	No
214	12/18	3	No
216	14/19	3	No
218	13/18	3	No
225	42/106	10	Yes
226	42/106	10	Elite
227	17/23	5	No
229	16/21	4	No
244	17/23	5	No
246	16/21	4	No
300	42/97	10	Elite
326	14/20	4	No
327	13/19	3	No
328	14/19	3	No
345	14/19	3	No
347	13/18	3	No
351	14/19	3	No
354	12/18	3	No
442	18/25	5	No
444	17/23	5	No
462	20/60	10	No
571	20/27	6	No
578	19/26	5	No
584	14/19	3	No
616	27/65	9	Yes
742	10/14	1	No

747	10/15	2	No
777	29/71	10	Elite
789	26/59	10	Elite
793	27/62	10	Elite

We're going to take each hybrid row and split them into two new rows - one with values for the first fuel type (values before the "/"), and the other with values for the second fuel type (values after the "/"). Let's separate them with two dataframes!

```
In [11]: # create two copies of the 2008 hybrids dataframe
df1 = hb_08.copy() # data on first fuel type of each hybrid vehicle
df2 = hb_08.copy() # data on second fuel type of each hybrid vehicle

# Each one should look like this
df1
```

```
Out[11]:
```

	model	displ	cyl	trans	drive	fuel	veh_class	\
582	MERCEDES-BENZ C300	3.0	6	Auto-L7	2WD	ethanol/gas	small car	

	air_pollution_score	city_mpg	hwy_mpg	cmb_mpg	greenhouse_gas_score	smartway
582	6/4	13/18	19/25	15/21		7/6 no

For this next part, we're going use Pandas' apply function. See the docs [here](#).

```
In [12]: # columns to split by "/"
split_columns = ['fuel', 'air_pollution_score', 'city_mpg', 'hwy_mpg', 'cmb_mpg', 'greenhouse_gas_score']

# apply split function to each column of each dataframe copy
for c in split_columns:
    df1[c] = df1[c].apply(lambda x: x.split("/")[0])
    df2[c] = df2[c].apply(lambda x: x.split("/")[1])
```

```
In [13]: # this dataframe holds info for the FIRST fuel type of the hybrid
# aka the values before the "/"s
df1
```

```
Out[13]:
```

	model	displ	cyl	trans	drive	fuel	veh_class	\
582	MERCEDES-BENZ C300	3.0	6	Auto-L7	2WD	ethanol	small car	

	air_pollution_score	city_mpg	hwy_mpg	cmb_mpg	greenhouse_gas_score	smartway
582	6	13	19	15		7 no

```
In [14]: # this dataframe holds info for the SECOND fuel type of the hybrid
# aka the values before the "/"s
df2
```

```
Out[14]:
```

	model	displ	cyl	trans	drive	fuel	veh_class	\
582	MERCEDES-BENZ C300	3.0	6	Auto-L7	2WD	gas	small car	

	air_pollution_score	city_mpg	hwy_mpg	cmb_mpg	greenhouse_gas_score	smartway
582	4	18	25	21		6 no

```

In [15]: # combine dataframes to add to the original dataframe
new_rows = df1.append(df2)

# now we have separate rows for each fuel type of each vehicle!
new_rows

Out[15]:
      model  displ  cyl  trans drive  fuel  veh_class \
582  MERCEDES-BENZ C300   3.0   6  Auto-L7   2WD  ethanol  small car
582  MERCEDES-BENZ C300   3.0   6  Auto-L7   2WD    gas  small car

      air_pollution_score  city_mpg  hwy_mpg  cmb_mpg  greenhouse_gas_score  smartway
582                      6         13        19        15                      7        no
582                      4         18        25        21                      6        no

In [16]: hb_08.index

Out[16]: Int64Index([582], dtype='int64')

In [17]: hb_18.index

Out[17]: Int64Index([ 52,  78,  79,  92, 189, 195, 196, 197, 212, 214, 216, 218, 225,
                    226, 227, 229, 244, 246, 300, 326, 327, 328, 345, 347, 351, 354,
                    442, 444, 462, 571, 578, 584, 616, 742, 747, 777, 789, 793],
                    dtype='int64')

In [18]: # drop the original hybrid rows
df_08.drop(hb_08.index, inplace=True)

# add in our newly separated rows
df_08 = df_08.append(new_rows, ignore_index=True)

In [19]: # check that all the original hybrid rows with "/"s are gone
df_08[df_08['fuel'].str.contains('/')]

Out[19]: Empty DataFrame
Columns: [model, displ, cyl, trans, drive, fuel, veh_class, air_pollution_score, city_mpg, hwy_mpg, cmb_mpg, greenhouse_gas_score, smartway]
Index: []

In [20]: df_08.shape

Out[20]: (987, 13)

```

4 Repeat this process for the 2018 dataset

```

In [21]: # create two copies of the 2018 hybrids dataframe, hb_18
df1 = hb_18.copy() # data on first fuel type of each hybrid vehicle
df2 = hb_18.copy()

```


4.0.1 Split values for fuel, city_mpg, hwy_mpg, cmb_mpg

You don't need to split for air_pollution_score or greenhouse_gas_score here because these columns are already ints in the 2018 dataset.

```
In [22]: # list of columns to split
        split_columns = ['fuel', 'city_mpg', 'hwy_mpg', 'cmb_mpg']

        # apply split function to each column of each dataframe copy
        for c in split_columns:
            df1[c] = df1[c].apply(lambda x: x.split("/")[0])
            df2[c] = df2[c].apply(lambda x: x.split("/")[1])
```

```
In [23]: df1.shape
```

```
Out[23]: (38, 13)
```

```
In [24]: df2.shape
```

```
Out[24]: (38, 13)
```

```
In [25]: # append the two dataframes
        new_rows = df1.append(df2)

        # drop each hybrid row from the original 2018 dataframe
        # do this by using Pandas drop function with hb_18's index
        df_18.drop(hb_18.index, inplace=True)

        # append new_rows to df_18
        df_18 = df_18.append(new_rows, ignore_index=True)
```

```
In [26]: # check that they're gone
        df_18[df_18['fuel'].str.contains('/')]
```

```
Out[26]: Empty DataFrame
        Columns: [model, displ, cyl, trans, drive, fuel, veh_class, air_pollution_score, city_mpg, hwy_mpg, cmb_mpg, greenhouse_gas_score]
        Index: []
```

```
In [27]: df_18.shape
```

```
Out[27]: (832, 13)
```

4.0.2 Now we can comfortably continue the changes needed for air_pollution_score! Here they are again:

- 2008: convert string to float
- 2018: convert int to float

```
In [28]: df_08['air_pollution_score'].value_counts()
```

```

Out[28]: 6      501
        7      398
        9.5    80
        9       7
        4       1
        Name: air_pollution_score, dtype: int64

In [29]: # convert string to float for 2008 air pollution column
        df_08['air_pollution_score'] = df_08['air_pollution_score'].astype(float)

In [30]: df_08['air_pollution_score'].value_counts()

Out[30]: 6.0      501
        7.0      398
        9.5      80
        9.0       7
        4.0       1
        Name: air_pollution_score, dtype: int64

In [31]: # convert int to float for 2018 air pollution column
        df_18['air_pollution_score'] = df_18['air_pollution_score'].astype(float)

In [32]: df_08['air_pollution_score'].value_counts()

Out[32]: 6.0      501
        7.0      398
        9.5      80
        9.0       7
        4.0       1
        Name: air_pollution_score, dtype: int64

In [33]: df_08.to_csv('data_08.csv', index=False)
        df_18.to_csv('data_18.csv', index=False)

In [ ]:

```